

What can a **Service Mesh** do for you?

Mike Stoltz



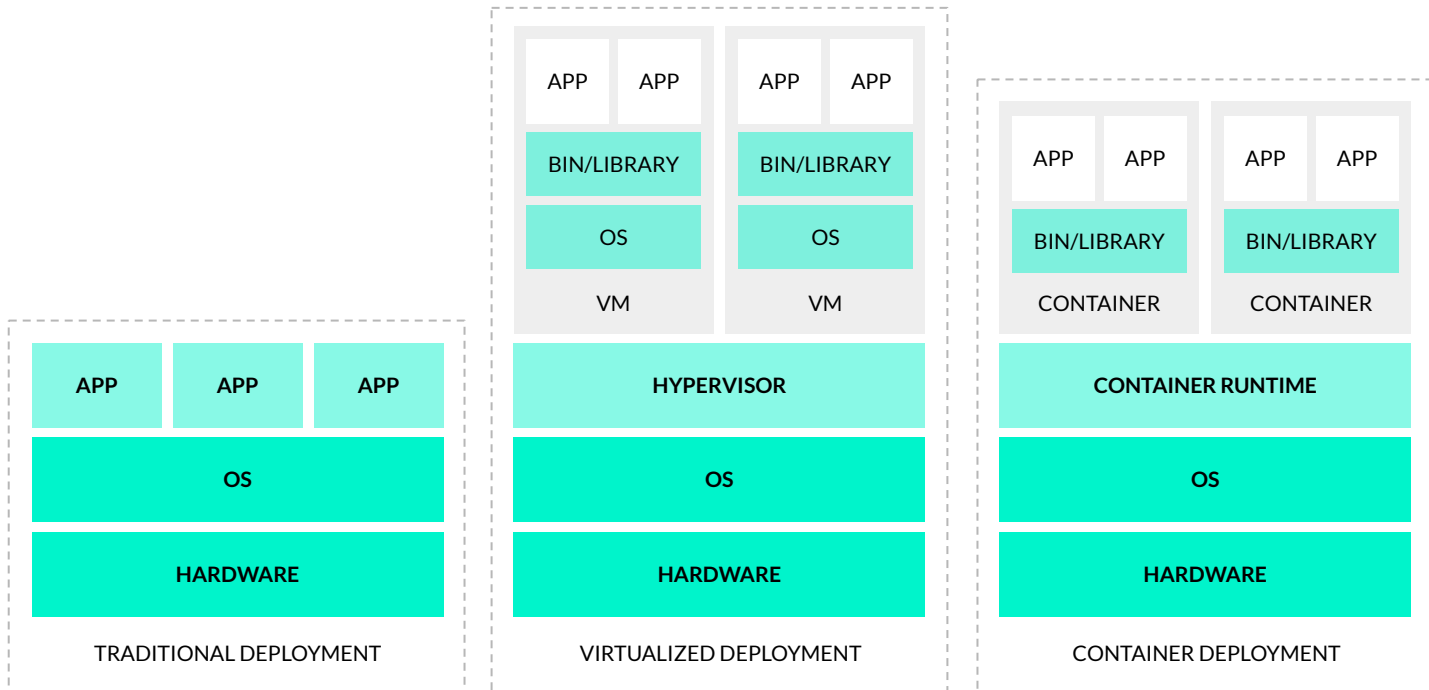
containership

Who am I?

- Mike Stoltz
- Software Engineer
- Containership Inc
- 4 Years Container Orchestrators
- 3 Years Kubernetes

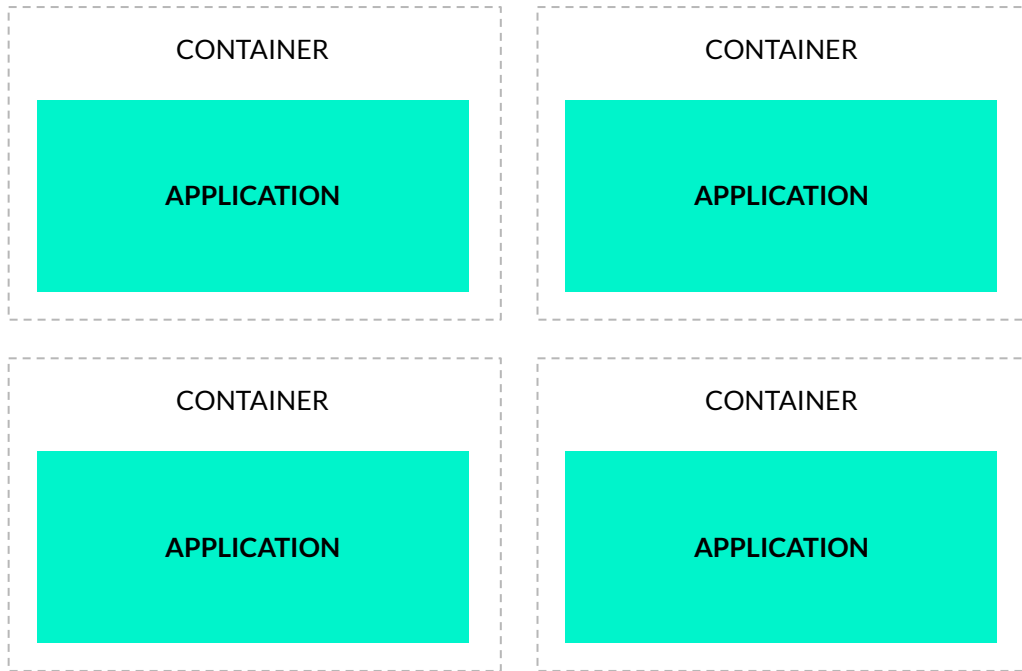


Virtualization History



Container Infrastructure Challenges

- Manual application scaling
- Networking & Routing
- Containers can run anywhere



Kubernetes

- Open source container orchestrator for automating application management and deployment across clusters of hosts.

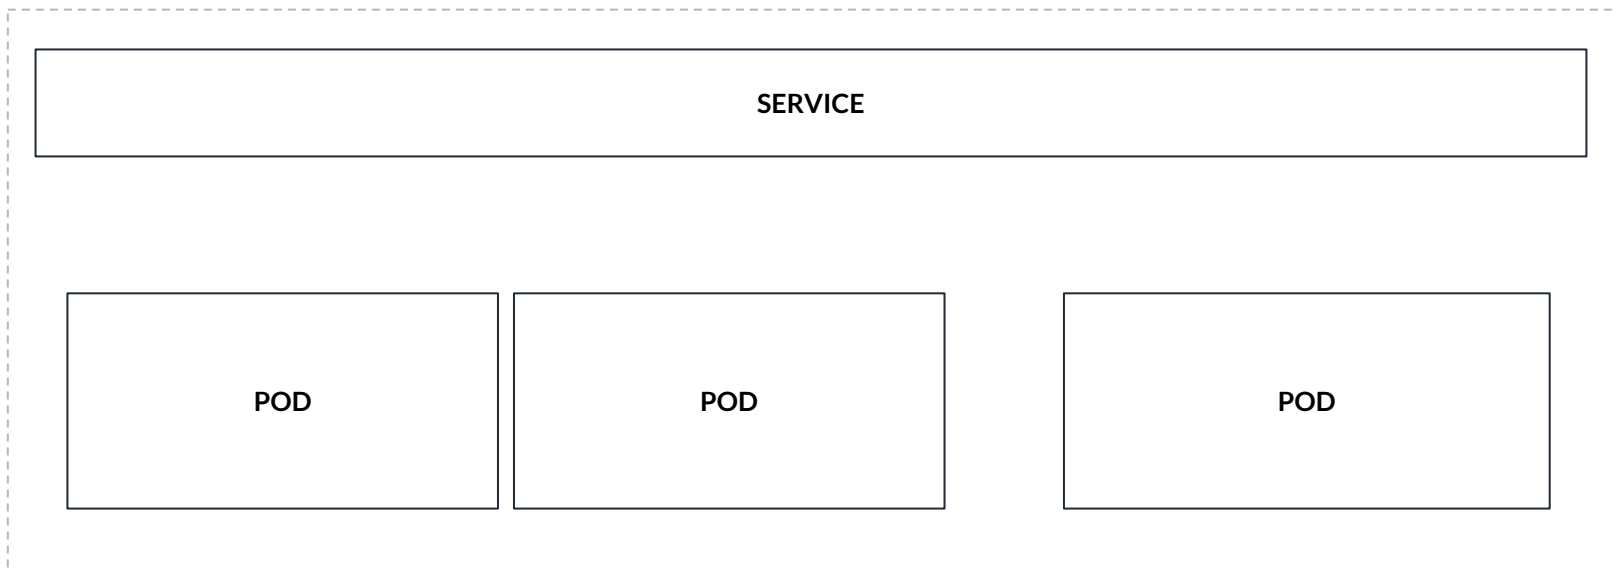


Kubernetes Terms

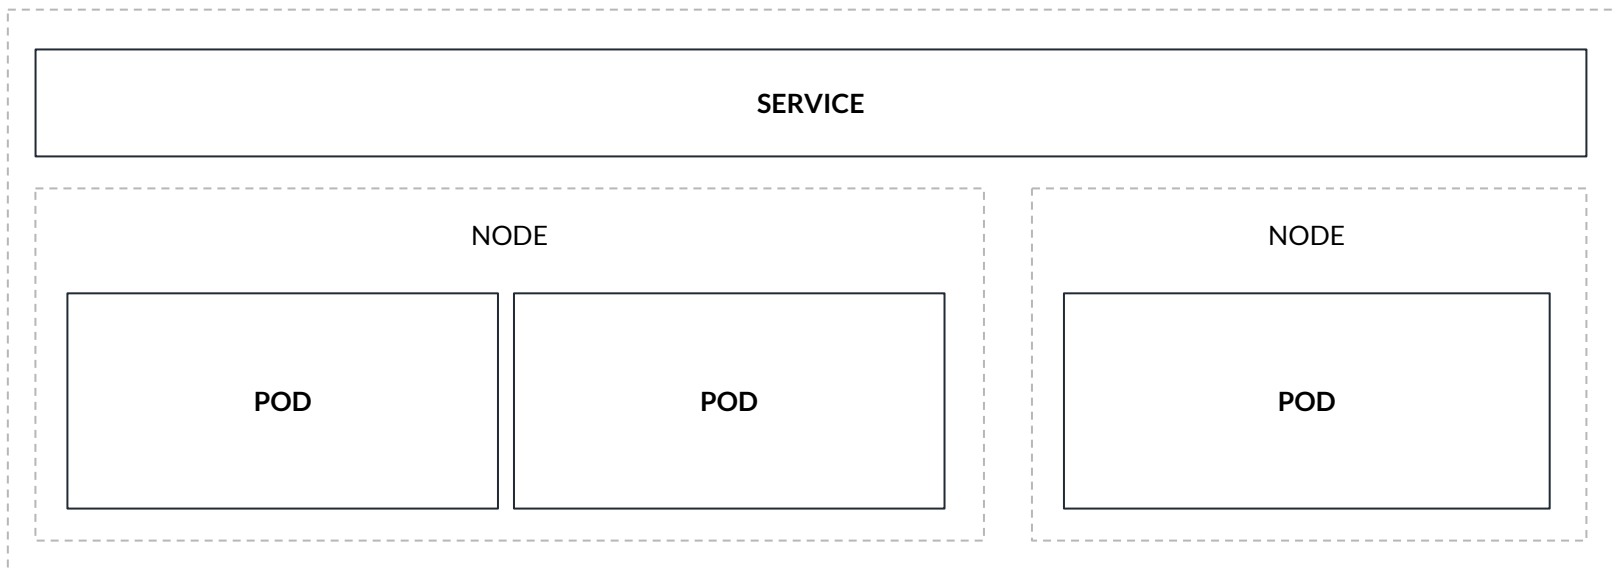
- **Node** - Virtual machine or physical server belonging to the cluster that containers will be scheduled on
- **Pod** - One or more containers and their environments that will share a virtual network
- **Service** - An abstraction which defines a logical set of Pods and a policy by which to access them
- **Selectors** - Labels used to match Kubernetes objects




Example Cluster




Example Cluster




Example Cluster



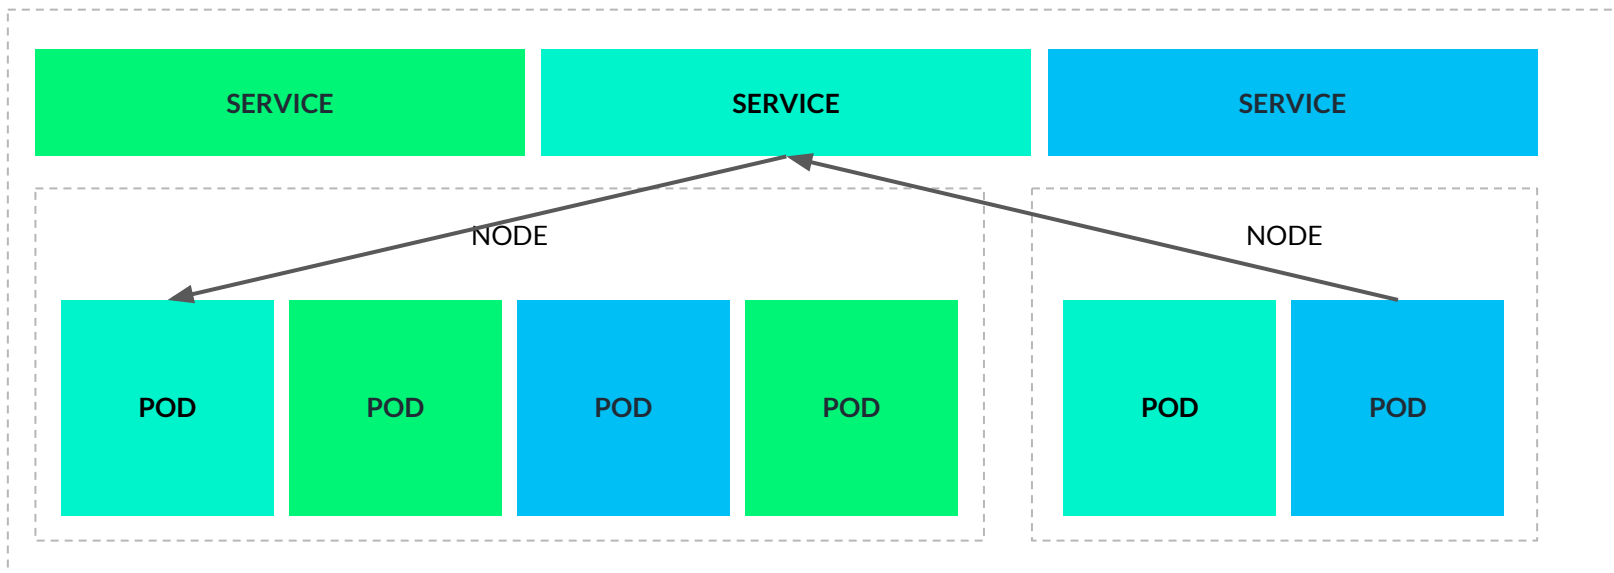
Example Cluster



Example Cluster



Example Cluster



Service Mesh - Advantages

- Infrastructure layer for handling service-to-service communication
- Observability
- Tracing
- Loadbalancing (L7)
- Encryption
- Authentication
- Separation of concerns

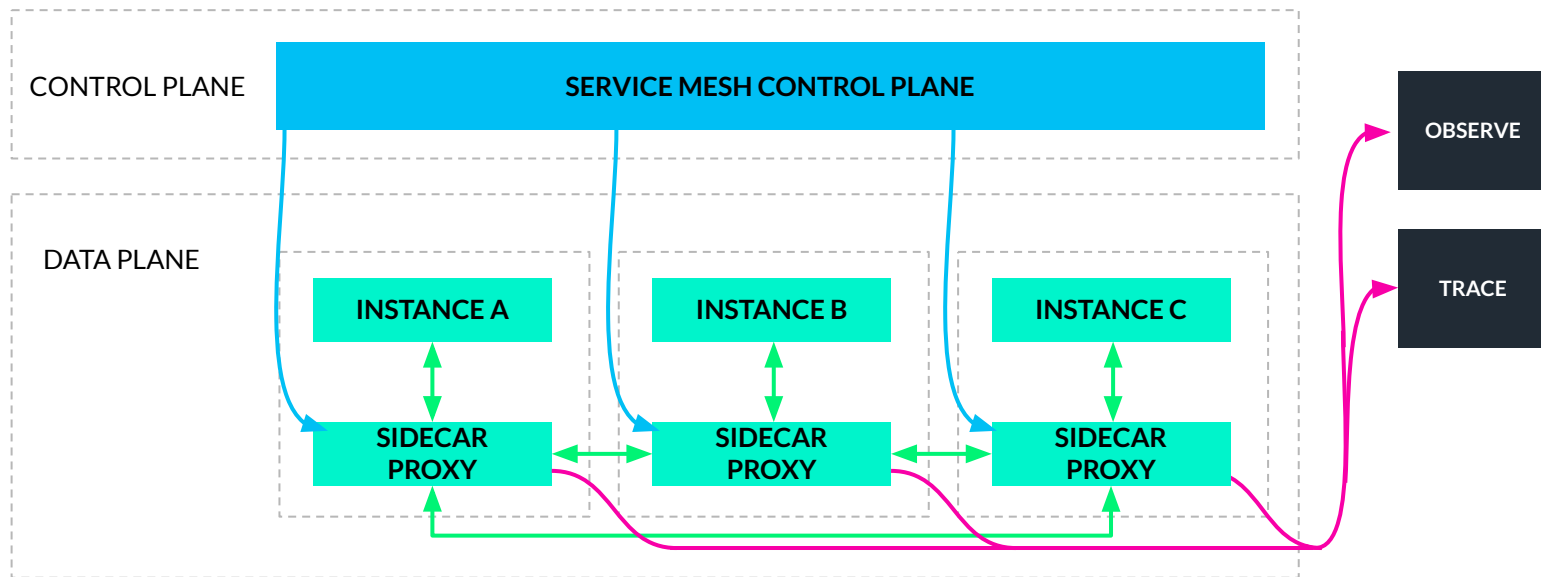


What makes up a service mesh?

- **Control plane** - a number of “services” running as controllers that accomplish things such as metrics collection and tracing, tls asset management, injecting the data plane proxy, and service discovery
- **Data plane** - a number of sidecar proxies that run alongside every pod, routes all traffic to and from the containers it is running alongside



What makes up a service mesh?



What service meshes are there?

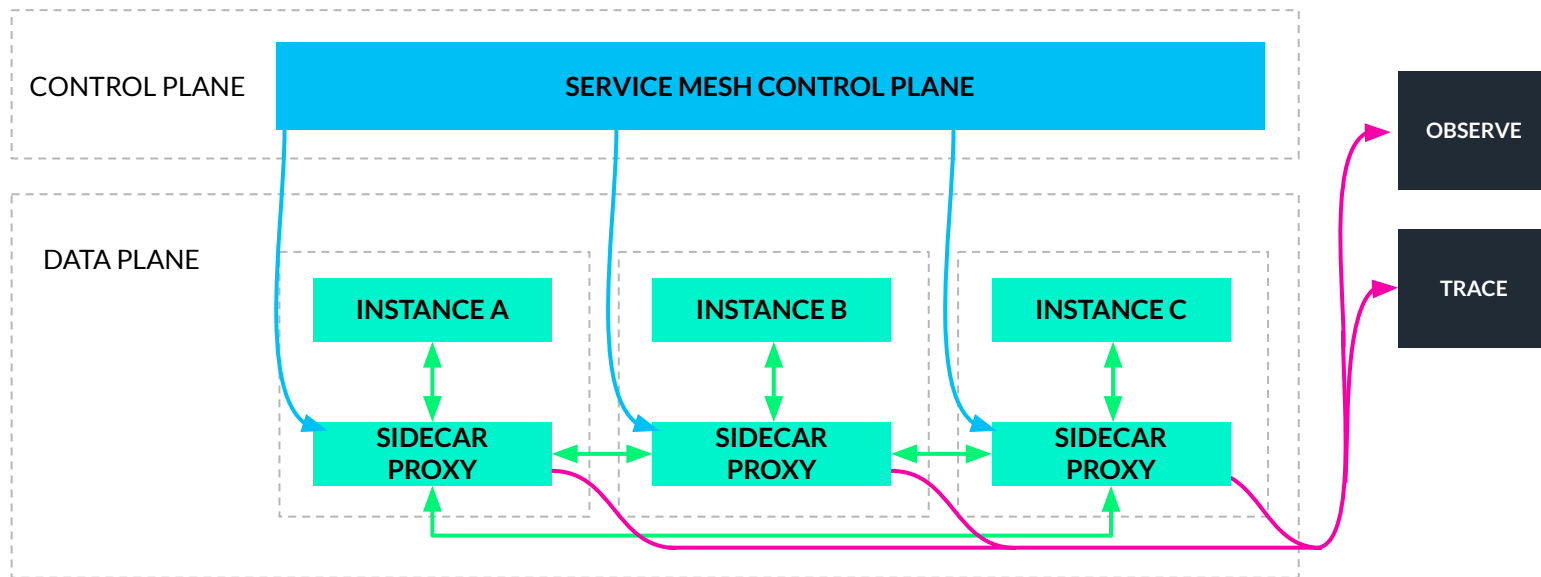


Linkerd



Istio

What makes up a service mesh?



Linkerd2 - Installation

```
$ curl -sL https://run.linkerd.io/install | sh
```

```
$ linkerd check --pre
```

```
$ linkerd install | kubectl apply -f -
```



Linkerd2 - Installation

```
$ namespace/linkerd created  
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-identity created  
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-identity created  
$ serviceaccount/linkerd-identity created  
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-controller created  
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-controller created  
$ serviceaccount/linkerd-controller created  
$ serviceaccount/linkerd-web created
```



Linkerd2 - Installation

```
$ customresourcedefinition.apiextensions.k8s.io/serviceprofiles.linkerd.io created
$ customresourcedefinition.apiextensions.k8s.io/trafficsplits.split.smi-spec.io created
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-prometheus created
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-prometheus created
$ serviceaccount/linkerd-prometheus created
$ serviceaccount/linkerd-grafana created
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-proxy-injector created
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-proxy-injector created
```



Linkerd2 - Installation

```
$ serviceaccount/linkerd-proxy-injector created
$ secret/linkerd-proxy-injector-tls created
$ mutatingwebhookconfiguration.admissionregistration.k8s.io/linkerd-proxy-injector-webhook created
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-sp-validator created
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-sp-validator created
$ serviceaccount/linkerd-sp-validator created
$ secret/linkerd-sp-validator-tls created
$ validatingwebhookconfiguration.admissionregistration.k8s.io/linkerd-sp-validator-webhook created
```



Linkerd2 - Installation

```
$ clusterrole.rbac.authorization.k8s.io/linkerd-linkerd-tap created
$ clusterrolebinding.rbac.authorization.k8s.io/linkerd-linkerd-tap created
$ serviceaccount/linkerd-tap created
$ podsecuritypolicy.policy/linkerd-linkerd-control-plane created
$ role.rbac.authorization.k8s.io/linkerd-psp created
$ rolebinding.rbac.authorization.k8s.io/linkerd-psp created
$ configmap/linkerd-config created
$ secret/linkerd-identity-issuer created
```

Linkerd2 - Installation

```
$ service/linkerd-identity created
$ deployment.extensions/linkerd-identity created
$ service/linkerd-controller-api created
$ service/linkerd-destination created
$ deployment.extensions/linkerd-controller created
$ service/linkerd-web created
$ deployment.extensions/linkerd-web created
$ configmap/linkerd-prometheus-config created
```



Linkerd2 - Installation

```
$ service/linkerd-prometheus created  
$ deployment.extensions/linkerd-prometheus created  
$ configmap/linkerd-grafana-config created  
$ service/linkerd-grafana created  
$ deployment.extensions/linkerd-grafana created  
$ deployment.apps/linkerd-proxy-injector created  
$ service/linkerd-proxy-injector created  
$ service/linkerd-sp-validator created
```



Linkerd2 - Installation

```
$ deployment.extensions/linkerd-sp-validator created  
$ service/linkerd-tap created  
$ deployment.extensions/linkerd-tap created
```

Linkerd2 - Installation

```
$ kubectl get pod -n linkerd
```

NAME	READY	STATUS	RESTARTS	AGE
Linkerd-controller-5bcf647776-qxxsd	3/3	Running	0	17m
Linkerd-grafana-64dc78cc97-52jdg	2/2	Running	0	16m
Linkerd-identity-7f84fcc96d-zqstr	2/2	Running	0	17m
Linkerd-prometheus-7df8689c4f-x64lr	2/2	Running	0	17m
Linkerd-proxy-injector-876d98c99-fb7fp	2/2	Running	0	16m
Linkerd-sp-validator-d6b4496d6-zjfv	2/2	Running	0	16m
linkerd-tap-7df59c45b7-ggbsv	2/2	Running	0	16m
linkerd-web-7ff6bbd984-fp5gc	2/2	Running	0	17m



Linkerd2 - Dashboard Demo

The screenshot shows the Linkerd2 dashboard for the 'linkerd' namespace. The sidebar on the left contains navigation links: Overview, Tap, Top, Top Routes, Service Mesh, Resources, Documentation, Community, Join the Mailing List, Join us on Slack, and File an Issue. The main content area displays the 'linkerd' namespace and a 'meshed' status. Below this, there are two tables: 'Deployments' and 'Pods'. Both tables show metrics for various components, including Success Rate, RPS, and Latency (P50, P95, P99).

Deployments

Deployment	Meshed	Success Rate	RPS	P50 Latency	P95 Latency	P99 Latency	Grafana
linkerd-controller	1/1	100.00%	1.9	240 ms	725 ms	945 ms	
linkerd-grafana	1/1	100.00%	0.3	1 ms	4 ms	4 ms	
linkerd-identity	1/1	100.00%	0.3	1 ms	2 ms	2 ms	
linkerd-prometheus	1/1	100.00%	27.43	171 ms	687 ms	937 ms	
linkerd-proxy-injector	1/1	100.00%	0.3	1 ms	9 ms	10 ms	
linkerd-sp-validator	1/1	100.00%	0.3	1 ms	1 ms	1 ms	
linkerd-tap	1/1	100.00%	0.3	1 ms	3 ms	3 ms	
linkerd-web	1/1	100.00%	0.3	1 ms	3 ms	3 ms	

Pods

Pod	Meshed	Success Rate	RPS	P50 Latency	P95 Latency	P99 Latency	Grafana
linkerd-controller-5bcf647776-qxxsd	1/1	100.00%	1.9	240 ms	725 ms	945 ms	
linkerd-grafana-64dc78cc97-52jdq	1/1	100.00%	0.3	1 ms	4 ms	4 ms	
linkerd-identity-7f8d6fcr-06d-rzotr	1/1	100.00%	0.3	1 ms	2 ms	2 ms	



Linkerd2

- Mutual TLS
- Automatic proxy injection
- L7 Loadbalancing

- Retries
- Timeouts
- Traffic splitting (not covered)



Linkerd2 - Service Profile

- Kubernetes CRD
- Provide extra information about a meshed service
- Routes
- Per route settings
- Retries
- Timeouts



Linkerd2 - Service Profile

```
spec:
  retryBudget:
    retryRatio: 0.2
    minRetriesPerSecond: 10
    ttl: 10s
  routes:
  - name: GET /api/annotations
    condition:
      method: GET
      pathRegex: /api/annotations
      isRetryable: true
```

```
spec:
  routes:
  - condition:
      method: HEAD
      pathRegex: /authors/[^/]*\.json
      name: HEAD /authors/{id}.json
      timeout: 300ms
```



Istio

- All covered Linkerd functionality
- Additional functionality
- Much more complex
- All new resource types
- Installed using Helm package manager



Istio - Custom Resources

- **Virtual Service** - a virtual service sets the default version and route logic for a microservice
- **Destination Rule** - defines policies that apply to traffic intended for a service after routing has occurred
- **Ingress Gateway** - describes a load balancer operating at the edge of the mesh receiving incoming or outgoing HTTP/TCP connections. The specification describes a set of ports that should be exposed, the type of protocol to use, SNI configuration for the load balancer
- **Service Entry** - used for stricter control of outbound requests to external resources, outside of the mesh and cluster



Istio - Virtual Service

```
spec:
  gateways:
    - bookinfo-gateway
  hosts:
    - details
  http:
    - match:
        - headers:
            end-user:
              exact: sam
      - route:
          - destination:
              host: details
              subset: v1
```

```
hosts:
  - reviews
http:
  - route:
      - destination:
          host: reviews
          subset: v1
        weight: 50
    - destination:
        host: reviews
        subset: v3
      weight: 50
    timeout: 0.5s
```



Downfalls

- Latency
- Resource overhead
- Proxy protocol support
- Increased complexity
- Application code changes for advanced tracing



Final Notes

- Kubernetes is not required for service meshes
- Observability is key
- Meshes add security
- Linkerd2 and Istio are both great solution
- Linkerd2 has small barrier to entry
- Istio allows for complex routing and settings



Thank you!

Mike Stoltz



@michael_stoltz



mike stoltz



containership